

MultiMarkdown v6 Quick Start Guide

Fletcher T. Penney

November 8, 2018

Contents

<i>Introduction</i>	1
<i>Performance</i>	1
<i>Parse Tree</i>	3
<i>Features</i>	3
<i>Abbreviations (Or Acronyms)</i>	3
<i>Citations</i>	4
<i>CriticMarkup</i>	4
<i>Embedded Images</i>	4
<i>Emph and Strong</i>	5
<i>EPUB 3 Support</i>	5
<i>Fenced Code Blocks</i>	5
<i>Footnotes</i>	5
<i>Glossary Terms</i>	6
<i>HTML Comments</i>	6
<i>Internationalization</i>	6
<i>LaTeX Changes</i>	6
<i>Metadata</i>	8
<i>Output Formats</i>	8
<i>Raw Source</i>	9
<i>Table of Contents</i>	10
<i>Tables</i>	10
<i>Transclusion</i>	10
<i>Developer Notes</i>	11
<i>Object Pools</i>	11
<i>HTML Boolean Attributes</i>	11
<i>Future Steps</i>	12

Introduction

Version: 6.4.0

This document serves as a description of MultiMarkdown (MMD) v6, as well as a sample document to demonstrate the various features. Specifically, differences from MMD v5 will be pointed out.

Performance

A big motivating factor leading to the development of MMD v6 was performance. When MMD first migrated from Perl to C (based on `peg-markdown`¹), it was among the fastest Markdown parsers available. That was many years ago, and the “competition” has made a great deal of progress since that time.

¹ <https://github.com/jgm/peg-markdown>

When developing MMD v6, one of my goals was to keep MMD at least in the ballpark of the fastest processors. Of course, being *the* fastest would be fantastic, but I was more concerned with ensuring that the code was easily understood, and easily updated with new features in the future.

MMD v3 – v5 used a PEG to handle the parsing. This made it easy to understand the relationship between the MMD grammar and the parsing code, since they were one and the same. However, the parsing code generated by the parsers was not particularly fast, and was prone to troublesome edge cases with terrible performance characteristics.

The first step in MMD v6 parsing is to break the source text into a series of tokens, which may consist of plain text, whitespace, or special characters such as “*”, “[”, etc. This chain of tokens is then used to perform the actual parsing.

MMD v6 divides the parsing into two separate phases, which actually fits more with Markdown’s design philosophically.

1. Block parsing consists of identifying the “type” of each line of the source text, and grouping the lines into blocks (e.g. paragraphs, lists, blockquotes, etc.) Some blocks are a single line (e.g. ATX headers), and others can be many lines long. The block parsing in MMD v6 is handled by a parser generated by `lemon`². This parser allows the block structure to be more readily understood by non-programmers, but the generated parser is still fast.
2. Span parsing consists of identifying Markdown/MMD structures that occur inside of blocks, such as links, images, strong, emph, etc. Most of these structures require matching pairs of tokens to specify where the span starts and where it ends. Most of these spans allow arbitrary levels of nesting as well. This made parsing

² <http://www.hwaci.com/sw/lemon/>

them correctly in the PEG-based code difficult and slow. MMD v6 uses a different approach that is accurate and has good performance characteristics even with edge cases. Basically, it keeps a stack of each “opening” token as it steps through the token chain. When a “closing” token is found, it is paired with the most recent appropriate opener on the stack. Any tokens in between the opener and closer are removed, as they are not able to be matched any more. To avoid unnecessary searches for non-existent openers, the parser keeps track of which opening tokens have been discovered. This allows the parser to continue moving forwards without having to go backwards and re-parse any previously visited tokens.

The result of this redesigned MMD parser is that it can parse short documents more quickly than CommonMark³, and takes only 15% – 20% longer to parse long documents. I have not delved too deeply into this, but I presume that CommonMark has a bit more “set-up” time that becomes expensive when parsing a short document (e.g. a paragraph or two). But this cost becomes negligible when parsing longer documents (e.g. file sizes of 1 MB). So depending on your use case, CommonMark may well be faster than MMD, but we’re talking about splitting hairs here. . . . Recent comparisons show MMD v6 taking approximately 4.37 seconds to parse a 108 MB file (approximately 24.8 MB/second), and CommonMark took 3.72 seconds for the same file (29.2 MB/second). For comparison, MMD v5.4 took approximately 94 second for the same file (1.15 MB/second).

³ <http://commonmark.org/>

For a more realistic file of approx 28 kb (the source of the Markdown Syntax web page), both MMD and CommonMark parse it too quickly to accurately measure. In fact, it requires a file consisting of the original file copied 32 times over (0.85 MB) before `/usr/bin/env time` reports a time over the minimum threshold of 0.01 seconds for either program.

There is still potentially room for additional optimization in MMD. However, even if I can’t close the performance gap with CommonMark on longer files, the additional features of MMD compared with Markdown in addition to the increased legibility of the source code of MMD (in my biased opinion anyway) make this project worthwhile.

Parse Tree

MMD v6 performs its parsing in the following steps:

1. Start with a null-terminated string of source text (C style string)
2. Lex string into token chain

3. Parse token chain into blocks
4. Parse tokens within each block into span level structures (e.g. strong, emph, etc.)
5. Export the token tree into the desired output format (e.g. HTML, LaTeX, etc.) and return the resulting C style string

OR

6. Use the resulting token tree for your own purposes.

The token tree (AST) includes starting offsets and length of each token, allowing you to use MMD as part of a syntax highlighter. MMD v5 did not have this functionality in the public version, in part because the PEG parsers used did not provide reliable offset positions, requiring a great deal of effort when I adapted MMD for use in MultiMarkdown Composer⁴.

⁴ <http://multimarkdown.com/>

These steps are managed using the `mmd_engine` “object”. An individual `mmd_engine` cannot be used by multiple threads simultaneously, so if `libMultiMarkdown` is to be used in a multithreaded program, a separate `mmd_engine` should be created for each thread. Alternatively, just use the slightly more abstracted `mmd_convert_string()` function that handles creating and destroying the `mmd_engine` automatically.

Features

Abbreviations (Or Acronyms)

This file includes the use of MMD as an abbreviation for MultiMarkdown. The abbreviation will be expanded on the first use, and the shortened form will be used on subsequent occurrences.

Abbreviations can be specified using inline or reference syntax. The inline variant requires that the abbreviation be wrapped in parentheses and immediately follows the `>`.

[>MMD] is an abbreviation. So is [>(MD) Markdown].

[>MMD]: MultiMarkdown

There is also a “shortcut” method for abbreviations that is similar to the approach used in prior versions of MMD. You specify the definition for the abbreviation in the usual manner, but MMD will automatically identify each instance where the abbreviation is used and substitute it automatically. In this case, the abbreviation is limited to a more basic character set which includes letters, numbers, periods, and hyphens, but not much else. For more complex abbreviations, you must explicitly mark uses of the abbreviation.

Citations

Citations can be specified using an inline syntax, just like inline footnotes. If you wish to use BibTeX, then configure the bibtex metadata (required) and the biblio style metadata (optional).

The HTML output for citations now uses parentheses instead of brackets, e.g. (1) instead of [1].

CriticMarkup

MMD v6 has improved support for CriticMarkup⁵, both in terms of parsing, and in terms of support for each output format. You can insert text, ~~delete text~~, substitute one thing for another, highlight text, and in the text.

⁵ <http://criticmarkup.com/>

leave comments

If you don't specify any command line options, then MMD will apply special formatting to the CriticMarkup formatting as in the preceding paragraph. Alternatively, you can use the `-a\`-accept or `-r\`-reject options to cause MMD to accept or reject, respectively, the proposed changes within the CM markup. When doing this, CM will work across blank lines. Without either of these two options, then CriticMarkup that spans a blank line is not recognized as such. I working on options for this for the future.

Embedded Images

Supported export formats (odt, epub, bundle, bundlezip) include images inside the export document:

- Local images are embedded automatically
- Images stored on remote servers are embedded *if* libCurl⁶ is properly installed when MMD is compiled. This is true for macOS builds.

⁶ <https://curl.haxx.se/libcurl/>

Emph and Strong

The basics of emphasis and strong emphasis are unchanged, but the parsing engine has been improved to be more accurate, particularly in various edge cases where proper parsing can be difficult.

EPUB 3 Support

MMD v6 now provides support for direct creation of EPUB 3⁷ files. Previously a separate tool was required to create EPUB files from MMD. It's now built-in. Currently, EPUB 3 files are built using the usual HTML 5 output. No extra CSS is applied, so the default

⁷ <https://en.wikipedia.org/wiki/EPUB>

from the reader will be used. Images are not yet supported, but are planned for the future.

EPUB files can be highly customized with other tools, and I recommend doing that for production quality files. For example, apparently performance is improved when the content is divided into multiple files (e.g. one file per chapter). MMD creates EPUB 3 files using a single file. Tools like Sigil⁸ are useful for improving your EPUB files, and I recommend doing that.

⁸ <https://sigil-ebook.com/>

Not all EPUB readers support v3 files. I don't plan on adding support for older versions of the EPUB format, but other tools can convert to other document formats you need. Same goes for Amazon's ebook formats – the Calibre⁹ program can also be used to interconvert between formats.

⁹ <https://calibre-ebook.com/>

NOTE: Because EPUB documents are binary files, MMD only creates them when run in batch mode (using the `-b\` `-batch` options). Otherwise, it simply outputs the HTML 5 file that would serve as the primary content for the EPUB.

Fenced Code Blocks

Fenced code blocks are fundamentally the same as MMD v5, except:

1. The leading and trailing fences can be 3, 4, or 5 backticks in length. That should be sufficient to account for complex documents without requiring a more complex parser.
2. If there is no trailing fence, then everything after the leading fence is considered to be part of the code block.

Footnotes

The HTML output for footnotes now uses superscripts instead of brackets, e.g. `¹` instead of `[1]`.

Glossary Terms

If there are terms in your document you wish to define in a glossary at the end of your document, you can define them using the glossary syntax.

Glossary terms can be specified using inline or reference syntax. The inline variant requires that the abbreviation be wrapped in parentheses and immediately follows the `?`.

`[?(glossary)` The glossary collects information about important terms used in your document] is a glossary term.

[?glossary] is also a glossary term.

[?glossary]: The glossary collects information about important terms used in your document

Much like abbreviations, there is also a “shortcut” method that is similar to the approach used in prior versions of MMD. You specify the definition for the glossary term in the usual manner, but MMD will automatically identify each instance where the term is used and substitute it automatically. In this case, the term is limited to a more basic character set which includes letters, numbers, periods, and hyphens, but not much else. For more complex glossary terms, you must explicitly mark uses of the term.

HTML Comments

Previously, HTML Comments were used by MultiMarkdown to include raw text for inclusion in the output file. This was useful, but limited, as it could only work for one output format at a time.

HTML Comments are now only included in HTML output, but not in any other format since they would cause errors.

Take a look at the `HTML Comments.text` file in the test suite for a better understanding of comment blocks vs comment spans, and how they are parsed.

Internationalization

MMD v6 includes support for substituting certain text phrases in other languages. This only affects the HTML format.

LaTeX Changes

LaTeX support is slightly different than in prior versions of MMD. It is designed to be a bit more consistent, and easier for basic use.

The previous approach used two types of metadata:

- `latex input` – this uses the name of a latex file that will be used in a `\input{file}` command. This key can be used multiple times (the only metadata key that worked this way), and all the basic metadata is written to the LaTeX file in order.
- `latex footer` – this file worked the same way as `latex input`, but was inserted at the end of the file

In practice, one typically needs to be able to insert `\input` commands at only a few key places in the final document:

1. At the very beginning
2. After metadata, and before the body of the document
3. After the body of the document

MMD 6 standardizes the metadata to use 3 new keys:

1. `latex leader` – this specifies a file that will be used at the very beginning of the document.
2. `latex begin` – this comes after metadata, and before the body of the document. This will usually include the `\begin{document}` command, hence the name.
3. `latex footer` – this comes after the body of the document.

You can use these 3 keys to replace the old `latex input` metadata keys, as long as you pay attention as to which is which. If you used more than three include statements, you may have to combine your latex files to fit into the new system.

In addition, there is a new shortcut key – `latex config`. This allows you to specify a “document name” that is used to automatically identify the corresponding `latex leader`, `latex begin`, and `latex footer` files. For example, using `latex config: article` is the same as using:

```
latex leader: mmd6-article-leader
latex begin: mmd6-article-begin
latex footer: mmd6-article-footer
```

Using the new system will require migrating your old configuration to the new naming convention, but once done I believe it should be much more intuitive to use.

The LaTeX support files included with the MMD v6 repository support the use of the following `latex config` values by default:

- `article`
- `tufte-book`
- `tufte-handout`

NOTE: You do have to install the MMD support files into the proper location for your system. I would like to make this easier, but haven’t found the best configuration yet.

Metadata

Metadata in MMD v6 includes new support for LaTeX – the `latex` config key allows you to automatically setup of multiple `latex` include files at once. The default setups that I use would typically consist of one LaTeX file to be included at the top of the file, one to be included right at the beginning of the document, and one to be included at the end of the document. If you want to specify the latex files separately, you can use `latex leader`, `latex begin`, and `latex footer`.

There are new metadata keys for controlling internationalization:

- `language` – specify the content language for a document, using the two letter code for the language (e.g. `en` for English). Where possible, this will also set the default `quotes language`.
- `quotes language` – specify which variant of smart quotes to use. Valid options are `dutch`, `french`, `german`, `germanguillemets`, `swedish`, `nl`, `fr`, `de`, `sv`. Anything else defaults to English.

Additionally, the `MMD Header` and `MMD Footer` metadata work slightly differently. In v5, these fields were used to list names of files that should be transcluded before and after the main body. In v6, these fields represent the actual text to be inserted. If you want them to reference separate files, use the transclusion functionality:

Title: Some Title

MMD Header: This is *MMD* text.

MMD Footer: {{footer.txt}}

Output Formats

MultiMarkdown 6 supports the following output formats, using the `-t` command-line argument:

- `html` – (Default) create HTML 5
- `latex` – create LaTeX¹⁰ for conversion to PDF using high quality typography
- `beamer` and `memoir` – two additional LaTeX variants for creating slide presentations and longer documents, respectively
- `mmd` – output the MMD text before converting to another format, but after performing transclusion. This format is not generally needed.
- `odt` – OpenDocument text file, used by OpenOffice and compatible word processors. Images are embedded inside the file package.

¹⁰ <https://en.wikipedia.org/wiki/LaTeX>

- `fodt` – OpenDocument text variant using a single text (XML) file instead of a compressed zip file. Images are not embedded in this format.
- `epub` – EPUB 3 ebook format. Images and CSS are embedded in the file package.
- `bundle` – [TextBundle] format consisting of Markdown/MultiMarkdown text file and embedded images and CSS. Useful for sharing Markdown files and images between applications (on any OS, but especially on iOS)
- `bundlezip` – TextPack variant of the TextBundle format – the file package is compressed to a single zip file (similar to EPUB and ODT formats).

Raw Source

In older versions of MultiMarkdown you could use an HTML comment to pass raw LaTeX or other content to the final document. This worked reasonably well, but was limited and didn't work well when exporting to multiple formats. It was time for something new.

MMD v6 offers a new feature to handle this. Code spans and code blocks can be flagged as representing raw source:

```
foo *bar*{=html}

'''{=latex}
*foo*
'''
```

The contents of the span/block will be passed through unchanged.

You can specify which output format is compatible with the specified source:

- `html`
- `odt`
- `epub`
- `latex`
- `*` – wildcard matches any output format

Table of Contents

By placing `{{T0C}}` in your document, you can insert an automatically generated Table of Contents in your document. As of MMD v6, the native Table of Contents functionality is used when exporting to LaTeX or OpenDocument formats.

Tables

Tables in MultiMarkdown-6 work basically the same as before, but a caption, if present, must come *after* the body of the table, not *before*.

Transclusion

File transclusion works basically the same way – `{{file}}` is used to indicate a file that needs to be transcluded. `{{file.*}}` allows for wildcard transclusion. What’s different is that the way search paths are handled is more flexible, though it may take a moment to understand.

When you process a file with MMD, it uses that file’s directory as the search path for included files. For example:

Directory	Transcluded Filename	Resolved Path
/foo/bar/	bat	/foo/bar/bat
/foo/bar/	baz/bat	/foo/bar/baz/bat
/foo/bar/	../bat	/foo/bat

This is the same as MMD v 5. What’s different is that when you transclude a file, the search path stays the same as the “parent” file, **UNLESS** you use the `transclude base` metadata to override it. The simplest override is:

```
transclude base: .
```

This means that any transclusions within the file will be calculated relative to the file, regardless of the original search path.

Alternatively you could specify that any transclusion happens inside a subfolder:

```
transclude base: folder/
```

Or you can specify an absolute path:

```
transclude base: /some/path
```

This flexibility means that you can transclude different files based on whether a file is being processed by itself or as part of a “parent” file. This can be useful when a particular file can either be a standalone document, or a chapter inside a larger document.

Developer Notes

If you’re using MMD as a library in another application, there are a few things to be aware of.

Object Pools

To improve performance, MMD has the option to allocate the memory for the tokens used in parsing in large chunks (“object pools”). Allocating a single large chunk of memory is more efficient than allocating many smaller chunks. However, this does complicate memory management.

By default `token.h` defines `kUseObjectPool` which enables this performance improvement. This does require more caution with the way that memory is managed. (See `main.c` for an example of how the object pool is allocated and drained.) I recommend disabling object pools unless you really understand C memory management, and understand MultiMarkdown’s program flow. Failure to properly manage the object pool can lead to massive memory leaks, freeing memory before that is still in use, or other potential problems.

HTML Boolean Attributes

Most HTML attributes are of the key-value type (e.g. `key="value"`). But some less frequently used attributes are boolean attributes (e.g. `<video controls>`). Properly distinguishing HTML from other uses of the `<` character requires matching both types under certain circumstances.

There are some trade-offs to be made:

- Performance when compiling MultiMarkdown
- Performance when processing parts of documents that are *not* HTML
- Accuracy when matching HTML

So far, there seem to be four main approaches:

- Ignore boolean attributes – this is how MMD-6 started. This is fast, but not accurate for some users. Several users found issues with the `<video>` tag when MMD was used in HTML heavy documents.
- Use regexp to match all boolean attributes. This is fast to compile, but adds roughly 5–8% processing time (probably due to false positive HTML matches). This *may* cause some text to be classified as HTML when it shouldn’t.
- Explicitly match all possible boolean attributes – This would presumably be relatively fast when processing (due to the nature of `re2c` lexers), but it may be prohibitively slow to compile for some users. As someone who compiles MMD frequently, it is too slow to compile be useful for me during development.

- Use a hand-curated list of boolean attributes that are most commonly used – this does not incur much of a performance hit when parsing, and compiles faster than the complete list of all boolean attributes. For now, this is the option I have chosen as default for MMD – it seems to be a reasonable trade-off. I will continue to research additional options.

Future Steps

Some features I plan to implement at some point:

1. OPML export support is not available in v6. I plan on adding improved support for this at some point. I was hoping to be able to re-use the existing v6 parser but it might be simpler to use the approach from v5 and earlier, which was to have a separate parser tuned to only identify headers and “stuff between headers”.

